

Navyatej Tummala

Raleigh, NC | (425) 515-1639 | navyatej31@gmail.com | [LinkedIn](#) | [GitHub](#)

EDUCATION

North Carolina State University <i>Master of Computer Science, GPA: 3.97</i>	Aug 2024 – May 2026 Raleigh, NC
GITAM University <i>Bachelor of Technology in Computer Science and Engineering, GPA: 7.93</i>	Aug 2020 – May 2024 Visakhapatnam, India

EXPERIENCE

Cloud Backend Engineer <i>ARoS Lab, NC State University, Raleigh, NC</i>	May 2025 – Aug 2025
<ul style="list-style-type: none">Engineered core features for a cloud web application to manage and label research image datasets, improving dataset readiness for downstream model training.Built a serverless backend using AWS Amplify and API Gateway, delivering REST APIs for dataset upload, labeling workflows, and export operations.Implemented event-driven ingestion with S3 and Lambda to ingest and preprocess large compressed image archives, improving reliability via trigger-based processing and structured logging.	
Software Engineering Intern <i>Symbiosis Technologies, Visakhapatnam, India</i>	May 2023 – May 2024

Software Engineering Researcher <i>GITAM University, Visakhapatnam, India</i>	Jan 2022 – May 2023
<ul style="list-style-type: none">Built a Python ETL with Pandas and NumPy to ingest and clean 5k+ news and social posts per day; exposed typed FastAPI endpoints with validation and caching, cutting preprocessing time about 40%.Built an OpenCV processing pipeline using Celery and Redis workers; persisted outputs to PostgreSQL with retries for resilient background execution.Delivered analytics APIs and React + Power BI dashboards, improving insight quality by about 20% and reducing refresh time by about 25%; productionized with Docker, GitHub Actions, pytest/mypy, auth, and rate limiting.	

PROJECTS

OnCall Copilot <i>Python, FAISS/Pinecone, LangChain, Slack Bolt, FastAPI, Docker, Redis</i>	
<ul style="list-style-type: none">Built an AI-powered Slack incident copilot that answers on-call questions from runbooks and postmortems, generating mitigation steps with linked dashboards and log queries.Implemented RAG with hybrid retrieval and reranking; deployed a containerized FastAPI service with Redis caching to achieve 89% top 3 source recall and about 780 ms median latency.	
GenAI Drug Information Assistant <i>Python, BERT, FDA API, Rasa, FastAPI, SQL</i>	
<ul style="list-style-type: none">Built a drug information chatbot using BERT-based NLP plus the FDA API to answer real-time medication and regulatory questions for clinicians and the public.Implemented the assistant in Rasa with intent classification, entity recognition, and dialogue management, plus a secure API and refresh pipeline to keep responses current.	

TECHNICAL SKILLS

Languages: Python, Java, TypeScript, SQL

Backend: FastAPI, Flask, REST APIs, OAuth2, JWT, OpenAPI

Cloud/DevOps: AWS (S3, Lambda, API Gateway, ECS/Fargate, CloudWatch), Docker

AI/ML: PyTorch, LangChain, scikit-learn, Sentence Transformers, RAG (FAISS/Pinecone)

Databases/Systems: PostgreSQL, MySQL, Redis, SQS, Celery, ETL pipelines

Testing/Quality: pytest, JUnit, mypy, observability (logging/metrics, OpenTelemetry)